


Introduction to Machine Learning with Scikit-Learn

Andreas Müller

Columbia University, scikit-learn

<https://github.com/amueller/ml-training-intro>





Lecture series:

<http://www.cs.columbia.edu/~amueller/comsw4995s18/schedule/>

What is machine learning?

Types of machine learning:

- supervised
- unsupervised
- reinforcement

Supervised Learning

$$(x_i, y_i) \propto p(x, y) \quad \text{i.i.d.}$$

$$x_i \in \mathbb{R}^n$$

$$y_i \in \mathbb{R}$$

$$f(x_i) \approx y_i$$

Classification and Regression

Classification:

- y discrete

Will you pass?

Regression:

- y continuous

How many points will you get in the exam?

Generalization

Not only

$$f(x_i) \approx y_i$$

Also for new data:

$$f(x) \approx y$$

Classification
Regression
Clustering
Semi-Supervised Learning
Feature Selection
Feature Extraction
Manifold Learning
Dimensionality Reduction
Kernel Approximation
Hyperparameter Optimization
Evaluation Metrics
Out-of-core learning

.....





Alexandre Gramfort
agramfort



Alexander Fabisch
AlexanderFabisch



Alexandre Passos
alextp



Andreas Mueller
amueller



Arnaud Joly
arjoly



Brian Holt
bdholt1



bthirion
bthirion



Chris Filo Gorgolewski
chrisfilo



David Cournapeau
cournape



Duchesnay
duchesnay



David Warde-Farley
dwf



Fabian Pedregosa
fabianp



Gael Varoquaux
GaelVaroquaux



Gilles Louppe
glouppe



Jake Vanderplas
jakevdp



Jaques Grobler
jaquesgrobler



Jan Hendrik Metzen
jmetzen



Jacob Schreiber
jmschrei



Joel Nothman
jnothman



Kyle Kastner
kastnerkyle



Lars
larsmans



Loïc Estève
lesteve



Shiqiao Du
lucidfrontier45



Mathieu Blondel
mblondel



Manoj Kumar
MechCoder



Noel Dawe
ndawe



Nelle Varoquaux
NelleV



Olivier Grisel
ogrisel



Paolo Losi
paolo-losi



Peter Prettenhofer
pprett



(Venkat) Raghav (Rajagopalan)
raghavrv



Robert Layton
robertlayton



Ron Weiss
ronw



Satrajit Ghosh
satra



sklearn-ci



sklearn-wheels



Tom Dupré la Tour
TomDLT



Vlad Niculae
vene



Virgile Fritsch
VirgileFritsch



Vincent Michel
vmichel





Wei Li
weillinear



Yaroslav Halchenko
yarikoptic

Get the notebooks!

nbviewer FAQ IPython  

pydata-nyc-advanced-sklearn / Chapter 0 - Reminder.ipynb /

Scikit-Learn is simple

Classification

```
In [4]: from sklearn.datasets import load_iris
        from sklearn.cross_validation import train_test_split

iris = load_iris()
X, y = iris.data, iris.target
X_train, X_test, y_train, y_test = train_test_split(X, y)

In [5]: from sklearn.svm import SVC
        clf = SVC()
        clf.fit(X_train, y_train)
        y_pred = clf.predict(X_test)
```

Transformations

```
In [6]: from sklearn.decomposition import PCA

In [7]: pca = PCA(n_components=2)
        pca.fit(X)
        X_pca = pca.transform(X)
```

Tools

Cross-validation scoring

```
In [30]: from sklearn.cross_validation import cross_val_score, StratifiedKFold
        scores = cross_val_score(SVC(), X_train, y_train, cv=5)
        print(scores)

[ 0.95652174  1.          0.95652174  0.91304348  0.9         ]
```

<https://github.com/amueller/ml-training-intro>

Documentation of scikit-learn 0.17

Quick Start

A very short introduction into machine learning problems and how to solve them using scikit-learn. Introduced basic concepts and conventions.

User Guide

The main documentation. This contains an in-depth description of all algorithms and how to apply them.

Other Versions

- [scikit-learn 0.18 \(development\)](#)
- [scikit-learn 0.17 \(stable\)](#)
- [scikit-learn 0.16](#)
- [scikit-learn 0.15](#)

Tutorials

Useful tutorials for developing a feel for some of scikit-learn's applications in the machine learning field.

API

The exact API of all functions and classes, as given by the docstrings. The API documents expected types and allowed features for all functions, and all parameters available for the algorithms.

Additional Resources

Talks given, slide-sets and other information relevant to scikit-learn.

Contributing

Information on how to contribute. This also contains useful information for advanced users, for example how to build their own estimators.

Flow Chart

A graphical overview of basic areas of machine learning, and guidance which kind of algorithms to use in a given situation.

FAQ

Frequently asked questions about the project and contributing.

Representing Data

one sample

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

one feature

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

outputs / labels

Training and Testing Data

training set

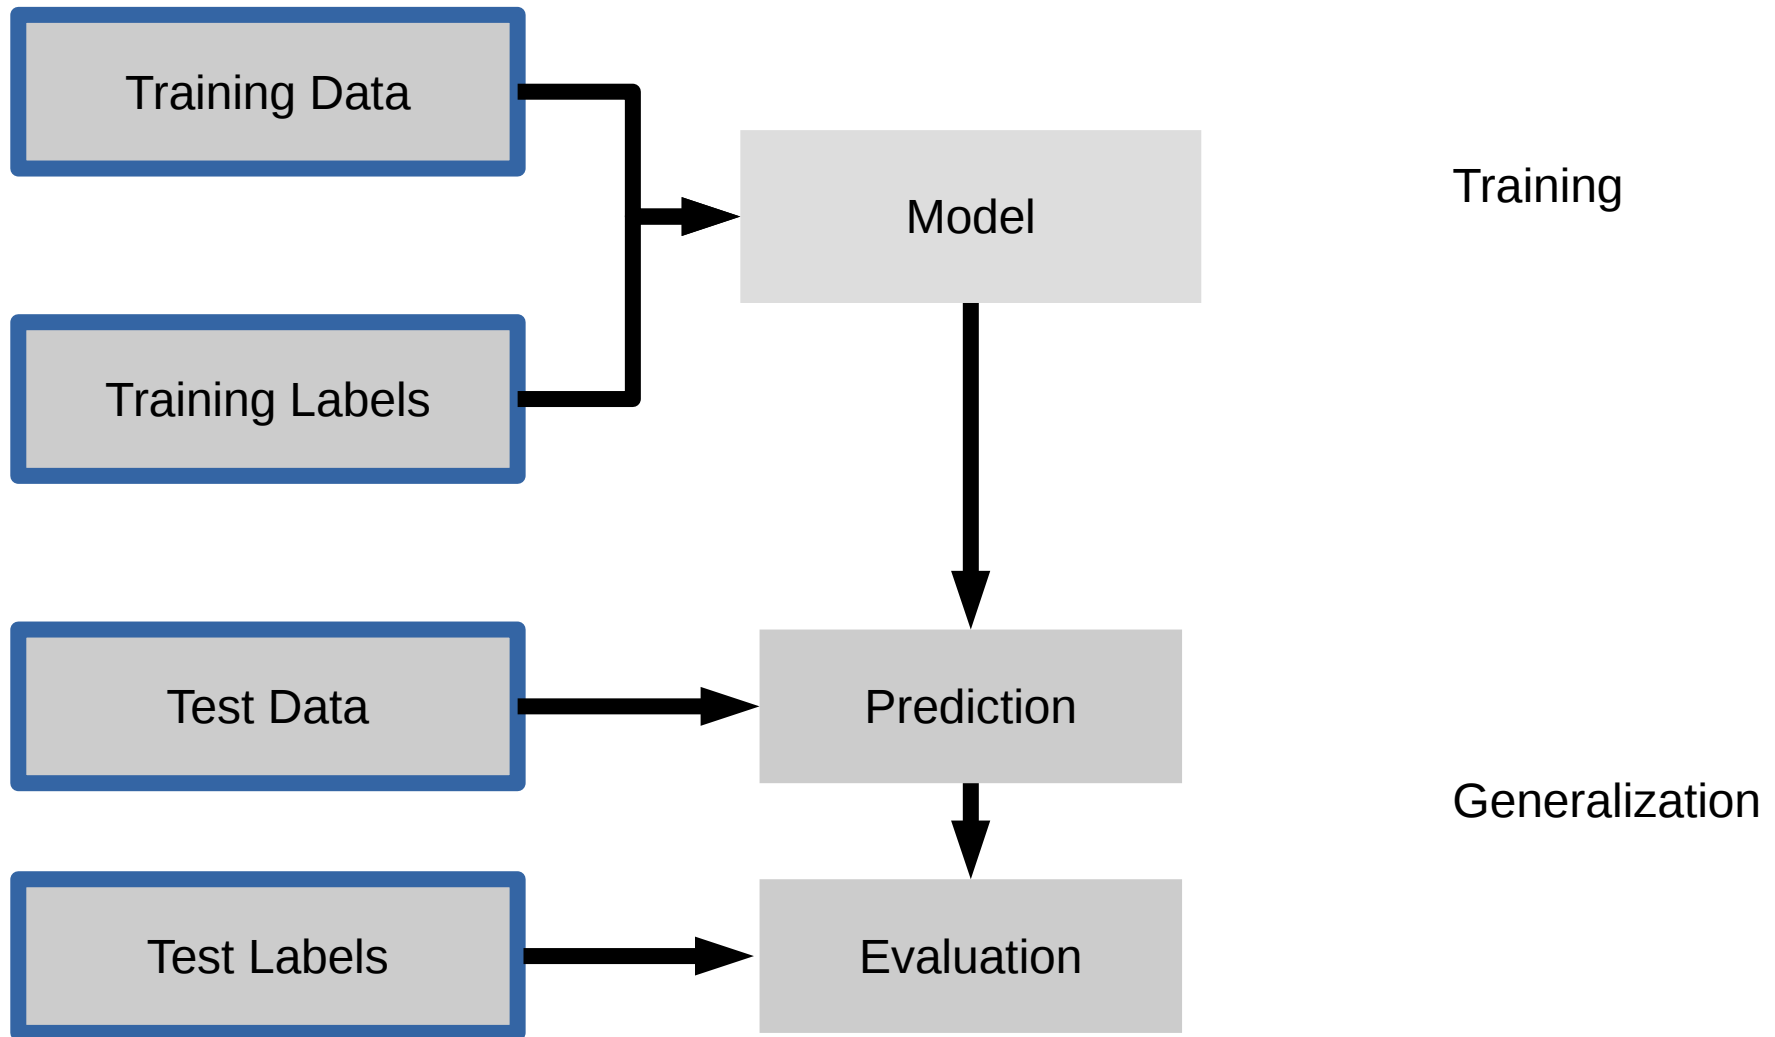
$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

test set

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

IPython Notebook: Part 0 – Data Loading

Supervised Machine Learning




```
clf = RandomForestClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```


```
clf.score(X_test, y_test)
```



IPython Notebook:

Part 1 - Introduction to Scikit-learn

Unsupervised Machine Learning




Unsupervised Transformations

```
pca = PCA()
```

```
pca.fit(X_train)
```

```
X_new = pca.transform(X_test)
```



IPython Notebook: Part 2 – Unsupervised Transformers

Basic API

`estimator.fit(X, [y])`

`estimator.predict`

`estimator.transform`

Classification

Preprocessing

Regression


Dimensionality reduction

Clustering

Feature selection


Feature extraction

Nearest neighbors







$$f(x) = y_i, i = \operatorname{argmin}_j ||x_j - x||$$


Nearest neighbors




Influence of $n_neighbors$




Model Complexity



Overfitting and Underfitting



Three-fold split



pro: fast, simple

con: high variance, bad use of data.

```
val_scores = []
neighbors = np.arange(1, 15, 2)
for i in neighbors:
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    val_scores.append(knn.score(X_val, y_val))
print("best validation score: {:.3f}".format(np.max(val_scores)))
best_n_neighbors = neighbors[np.argmax(val_scores)]
print("best n_neighbors: {}".format(best_n_neighbors))


knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(X_trainval, y_trainval)
print("test-set score: {:.3f}".format(knn.score(X_test, y_test)))
```

best validation score: 0.972

best n_neighbors: 3

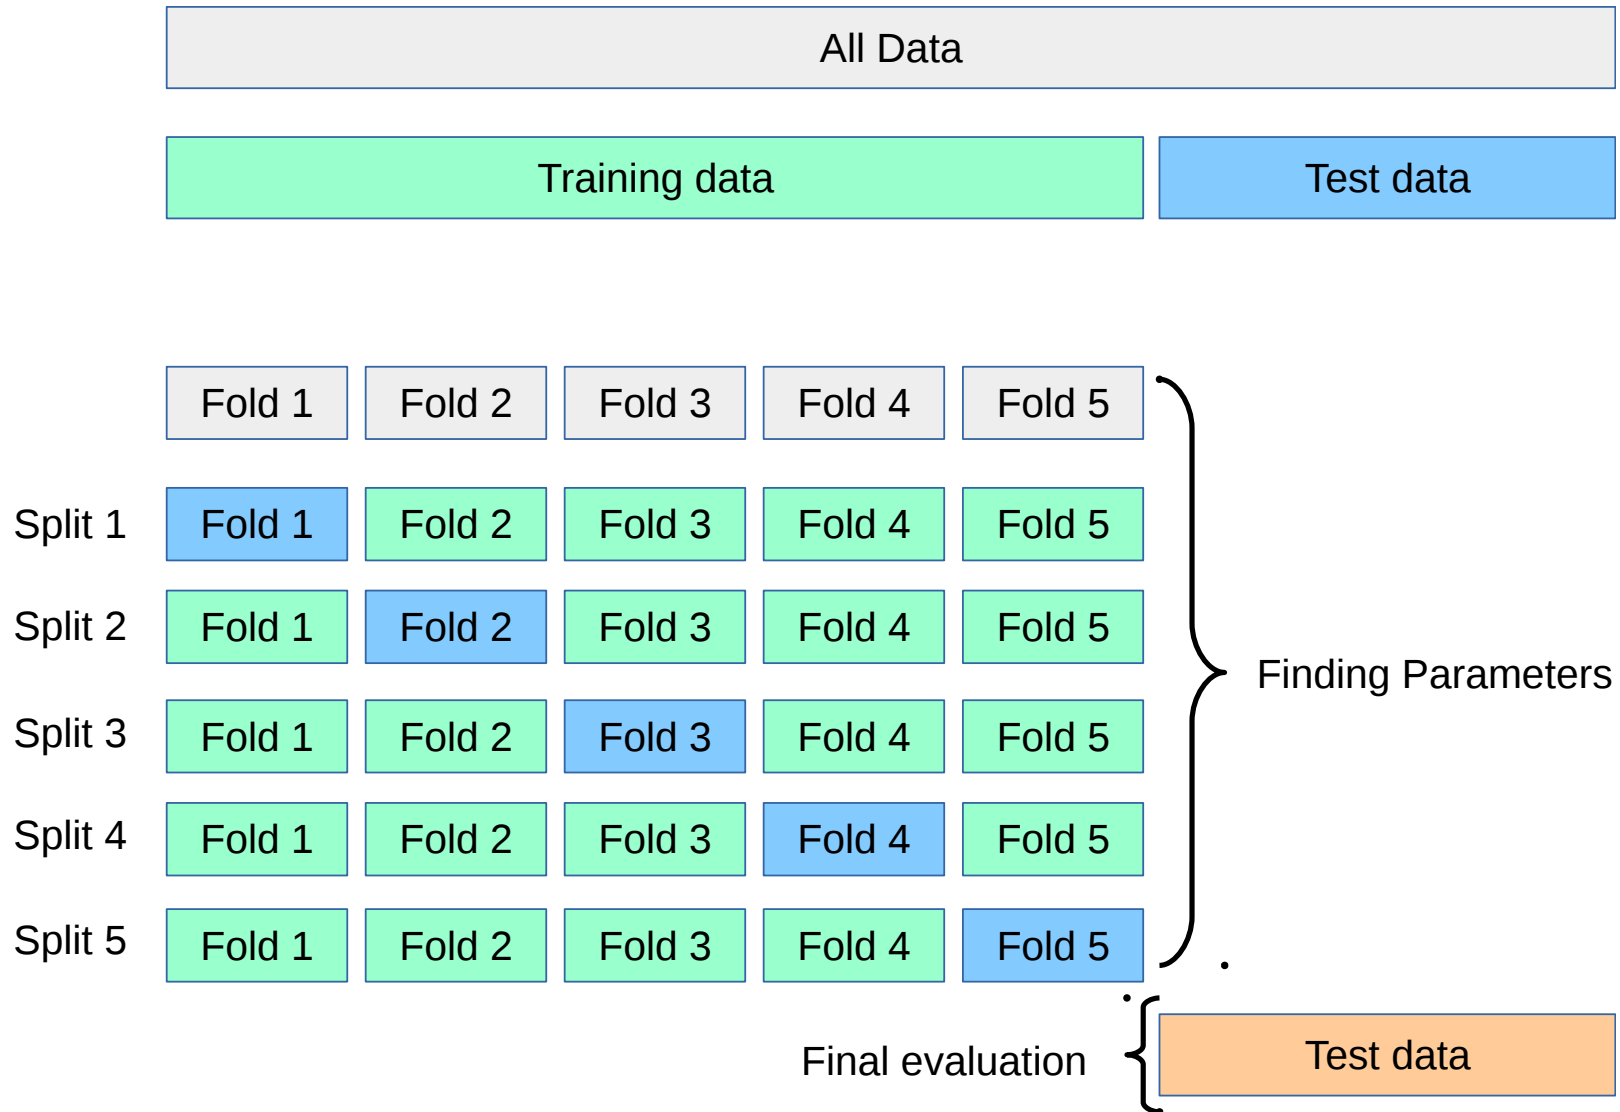
test-set score: 0.965

Cross-validation



Pro: more stable, more data
con: slower

Cross-validation + test-set



```
from sklearn.model_selection import cross_val_score

X_train, X_test, y_train, y_test = train_test_split(X, y)


cross_val_scores = []

for i in neighbors:
    knn = KNeighborsClassifier(n_neighbors=i)
    scores = cross_val_score(knn, X_trainval, y_trainval, cv=10)
    cross_val_scores.append(np.mean(scores))

print("best cross-validation score: {:.3f}".format(np.max(cross_val_scores)))
best_n_neighbors = neighbors[np.argmax(cross_val_scores)]
print("best n_neighbors: {}".format(best_n_neighbors))

knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(X_train, y_train)
print("test-set score: {:.3f}".format(knn.score(X_test, y_test)))
```

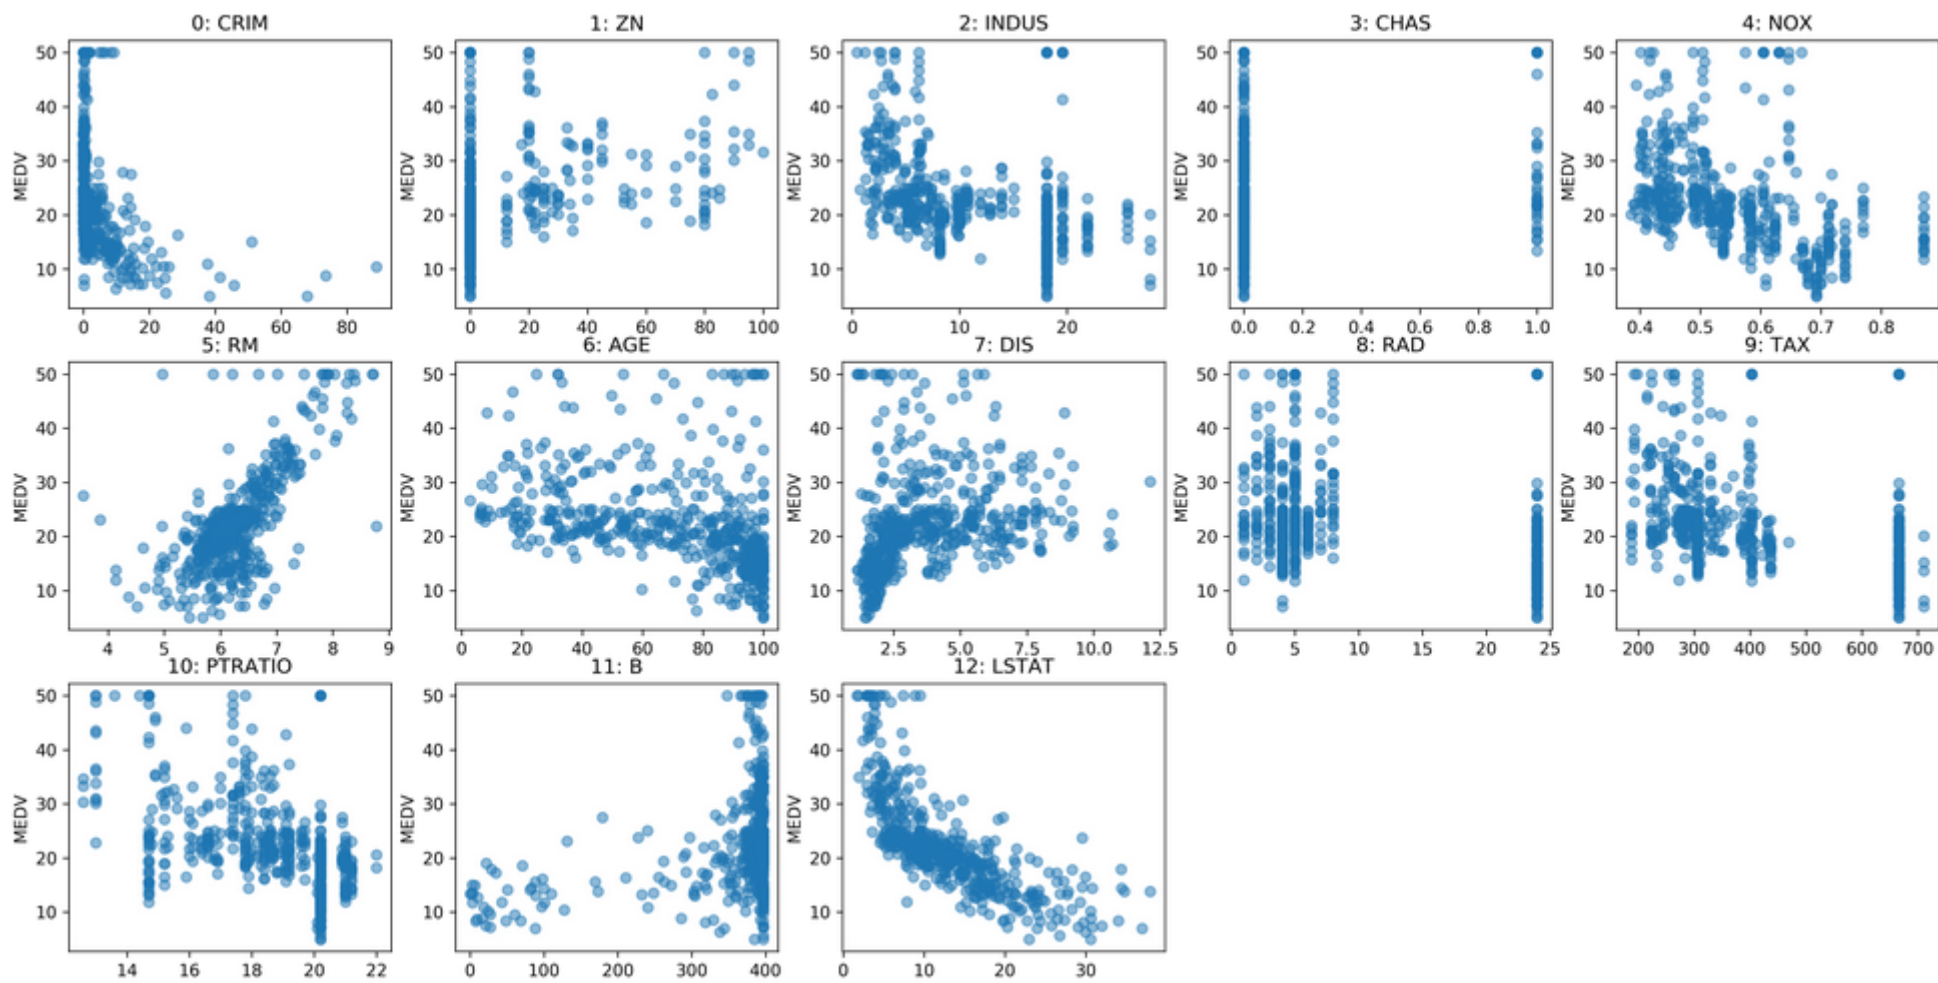
```
best cross-validation score: 0.972
best n_neighbors: 3
test-set score: 0.972
```




IPython Notebook:

Part 3 – Cross-validation and grid-search

Preprocessing



```
: plt.boxplot(X)
plt.xticks(np.arange(1, X.shape[1] + 1), boston.feature_names, rotation=30, ha="right")
plt.ylabel("MEDV")
: <matplotlib.text.Text at 0x7f580303eac8>
```






```
from sklearn.linear_model import Ridge
X, y = boston.data, boston.target
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
ridge = Ridge().fit(X_train_scaled, y_train)

X_test_scaled = scaler.transform(X_test)
ridge.score(X_test_scaled, y_test)
```

0.63448846877867426



Categorical Features

Categorical Features

$\{\text{"red"}, \text{"green"}, \text{"blue"}\} \subset \mathbb{R}^p$?


Categorical Variables

	“red”	“green”	“blue”
	1	0	0
	0	1	0
	0	0	1

IPython Notebook: Part 4 – Preprocessing

Linear Models for Regression

Linear Models for Regression



$$\hat{y} = w^T \mathbf{x} + b = \sum_{i=1}^p w_i x_i + b$$

Linear Regression

Ordinary Least Squares

$$\hat{y} = w^T \mathbf{x} + b = \sum_{i=1}^p w_i x_i + b$$

$$\min_{w \in \mathbb{R}^p} \sum_{i=1}^p ||w^T \mathbf{x}_i - y_i||^2$$

Unique solution if $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T$ has full rank.

Ridge Regression

$$\min_{w \in \mathbb{R}^p} \sum_{i=1}^n ||w^T x_i - y_i||^2 + \alpha ||w||^2$$


Always has a unique solution.

Has tuning parameter alpha

IPython Notebook:
Part 5 – Linear Models for Regression

Linear Models for Classification

Linear models for **binary** classification




$$\hat{y} = \text{sign}(w^T \mathbf{x} + b) = \text{sign}\left(\sum_i w_i x_i + b\right)$$

Logistic Regression

$$\min_{w \in \mathbb{R}^p} - \sum_{i=1}^n \log(\exp(-y_i w^T \mathbf{x}_i) + 1)$$

$$p(y|\mathbf{x}) = \frac{1}{1 + e^{-w^T \mathbf{x}}}$$




$$\hat{y} = \text{sign}(w^T \mathbf{x} + b)$$


Penalized Logistic Regression

$$\min_{w \in \mathbb{R}^p} -C \sum_{i=1}^n \log(\exp(-y_i w^T \mathbf{x}_i) + 1) + ||w||_2^2$$


C is inverse to alpha (or alpha / n_samples)




$C = 0.010000$



$C = 10.000000$



$C = 1000.000000$




Multinomial Logistic Regression

Probabilistic multi-class model:

$$p(y = i | \mathbf{x}) = \frac{e^{-\mathbf{w}_i^T \mathbf{x}}}{\sum_{j \in Y} e^{-\mathbf{w}_j^T \mathbf{x}}}$$

$$\min_{w \in \mathbb{R}^p} - \sum_{i=1}^n \log(p(y = y_i | \mathbf{x}_i))$$

$$\hat{y} = \arg \max_{i \in Y} \mathbf{w}_i \mathbf{x}$$

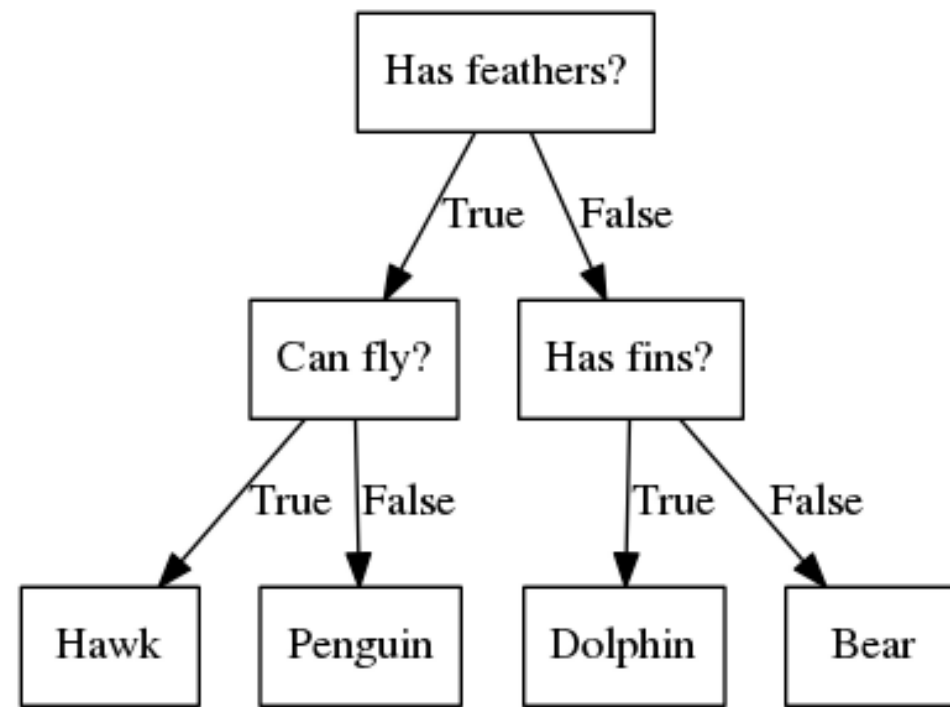


$$\hat{y} = \arg \max_{i \in Y} \mathbf{w}_i \mathbf{x}$$

IPython Notebook: Part 6 – Linear Models for Classification



Decision Trees and Tree-based Models

Idea: series of binary questions





Building trees

depth = 1



depth = 2



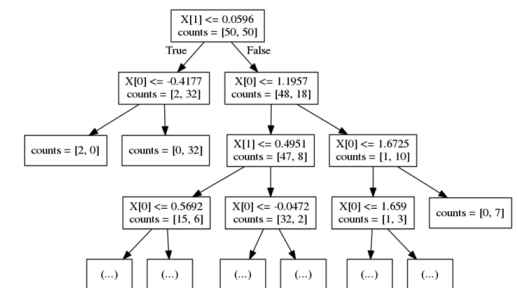

Continuous features:
“questions” are thresholds on
single features.

[Other methods are possible
but not as common]



For each split:
exhaustive search over all
features and thresholds!

Minimize “impurity”

depth = 9



Prediction




Traverse tree based on feature tests
Predict most common class in leaf


IPython Notebook: Part 7 – Decision Trees

Random Forests


tree 0




tree 1




tree 2




tree 3



tree 4




random forest




Randomize in two ways

- For each **tree**:
Pick bootstrap sample of data



- For each **split**:
Pick random sample of features




- More tree are always better

Tuning Random Forests

- Main parameter: `max_features`
 - around $\sqrt{n_features}$ for classification
 - Around `n_features` for regression
- `n_estimators > 100`
- Prepruning might help, definitely helps with model size!
- `max_depth`, `max_leaf_nodes`, `min_samples_split` again

O'REILLY®



Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido




amueller.github.io



[@amuellerml](https://twitter.com/amuellerml)



[@amueller](https://github.com/amueller)



andreas.mueller@columbia.edu