

# Introduction to Machine Learning with Scikit-Learn

Andreas Müller

Columbia University, scikit-learn

<https://github.com/amueller/ml-training-intro>



What is machine learning?

Types of machine learning:

- supervised
- unsupervised
- reinforcement

# Supervised Learning

$$(x_i, y_i) \propto p(x, y) \quad \text{i.i.d.}$$

$$x_i \in \mathbb{R}^n$$

$$y_i \in \mathbb{R}$$

$$f(x_i) \approx y_i$$

# Classification and Regression

Classification:

- $y$  discrete

Will you pass?

Regression:

- $y$  continuous

How many points will you get in the exam?

# Generalization

Not only


$$f(x_i) \approx y_i$$

Also for new data:

$$f(x) \approx y$$

Classification  
Regression  
Clustering  
Semi-Supervised Learning  
Feature Selection  
Feature Extraction  
Manifold Learning  
Dimensionality Reduction  
Kernel Approximation  
Hyperparameter Optimization  
Evaluation Metrics  
Out-of-core learning  
.....





Alexandre Gramfort  
agramfort



Alexander Fabisch  
AlexanderFabisch



Alexandre Passos  
alextp



Andreas Mueller  
amueller



Arnaud Joly  
arjoly



Brian Holt  
bdholt1



bthirlion  
bthirlion



Chris Filo Gorgolewski  
chrisfilo



David Courneau  
cournape



Duchesnay  
duchesnay



David Warde-Farley  
dwf



Fabian Pedregosa  
fabianp



Gael Varoquaux  
GaelVaroquaux



Gilles Louppe  
glouppe



Jake Vanderplas  
jakevdp



Jaques Grobler  
jaquesgrobler



Jan Hendrik Metzen  
jmetzen



Jacob Schreiber  
jmschrei



Joel Nothman  
jnothman



Kyle Kastner  
kastnerkyle



Lars  
larsmans



Loïc Estève  
lesteve



Shiqiao Du  
lucidfrontier45



Mathieu Blondel  
mblondel



Manoj Kumar  
MechCoder



Noel Dawe  
ndawe



Nelle Varoquaux  
NelleV



Olivier Grisel  
ogrisel



Paolo Losi  
paolo-losi



Peter Prettenhofer  
pprett



(Venkat) Raghav (Rajagopalan)  
raghavrv



Robert Layton  
robertlayton



Ron Weiss  
ronw



Satrajit Ghosh  
satra



sklearn-ci



sklearn-wheels



Tom Dupré la Tour  
TomDLT



Vlad Niculae  
vene



Virgile Fritsch  
VirgileFritsch



Vincent Michel  
vmichel



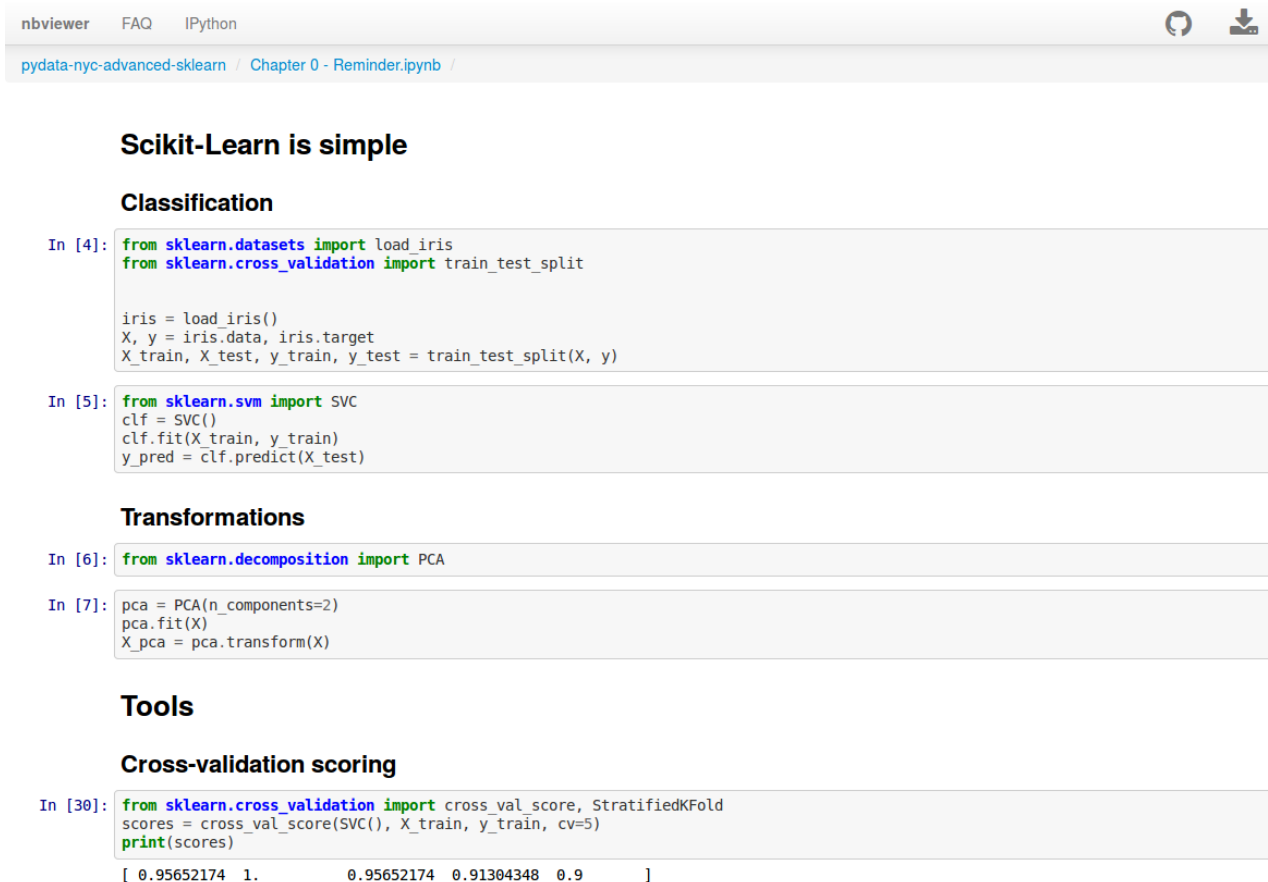
Wei Li  
weilinear



Yaroslav Halchenko  
yarikoptic



# Get the notebooks!



The screenshot shows a Jupyter Notebook interface with a top bar containing 'nbviewer', 'FAQ', and 'IPython' links, along with refresh and download icons. The breadcrumb path is 'pydata-nyc-advanced-sklearn / Chapter 0 - Reminder.ipynb /'. The notebook content is organized into sections: 'Scikit-Learn is simple', 'Classification', 'Transformations', 'Tools', and 'Cross-validation scoring'. Each section contains one or more code cells with Python code for loading data, training models, and evaluating performance.

```
nbviewer  FAQ  IPython
pydata-nyc-advanced-sklearn / Chapter 0 - Reminder.ipynb /

Scikit-Learn is simple

Classification

In [4]: from sklearn.datasets import load_iris
        from sklearn.cross_validation import train_test_split

        iris = load_iris()
        X, y = iris.data, iris.target
        X_train, X_test, y_train, y_test = train_test_split(X, y)

In [5]: from sklearn.svm import SVC
        clf = SVC()
        clf.fit(X_train, y_train)
        y_pred = clf.predict(X_test)

Transformations

In [6]: from sklearn.decomposition import PCA

In [7]: pca = PCA(n_components=2)
        pca.fit(X)
        X_pca = pca.transform(X)

Tools

Cross-validation scoring

In [30]: from sklearn.cross_validation import cross_val_score, StratifiedKFold
        scores = cross_val_score(SVC(), X_train, y_train, cv=5)
        print(scores)

[ 0.95652174  1.          0.95652174  0.91304348  0.9          ]
```

<https://github.com/amueller/ml-training-intro>

## Documentation of scikit-learn 0.17

### Quick Start

A very short introduction into machine learning problems and how to solve them using scikit-learn. Introduced basic concepts and conventions.

### User Guide

The main documentation. This contains an in-depth description of all algorithms and how to apply them.

### Other Versions

- [scikit-learn 0.18 \(development\)](#)
- [scikit-learn 0.17 \(stable\)](#)
- [scikit-learn 0.16](#)
- [scikit-learn 0.15](#)

### Tutorials

Useful tutorials for developing a feel for some of scikit-learn's applications in the machine learning field.

### API

The exact API of all functions and classes, as given by the docstrings. The API documents expected types and allowed features for all functions, and all parameters available for the algorithms.

### Additional Resources

Talks given, slide-sets and other information relevant to scikit-learn.

### Contributing

Information on how to contribute. This also contains useful information for advanced users, for example how to build their own estimators.

### Flow Chart

A graphical overview of basic areas of machine learning, and guidance which kind of algorithms to use in a given situation.

### FAQ

Frequently asked questions about the project and contributing.

<http://scikit-learn.org/>

# Representing Data

one sample

$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

one feature

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

outputs / labels

# Training and Testing Data

training set

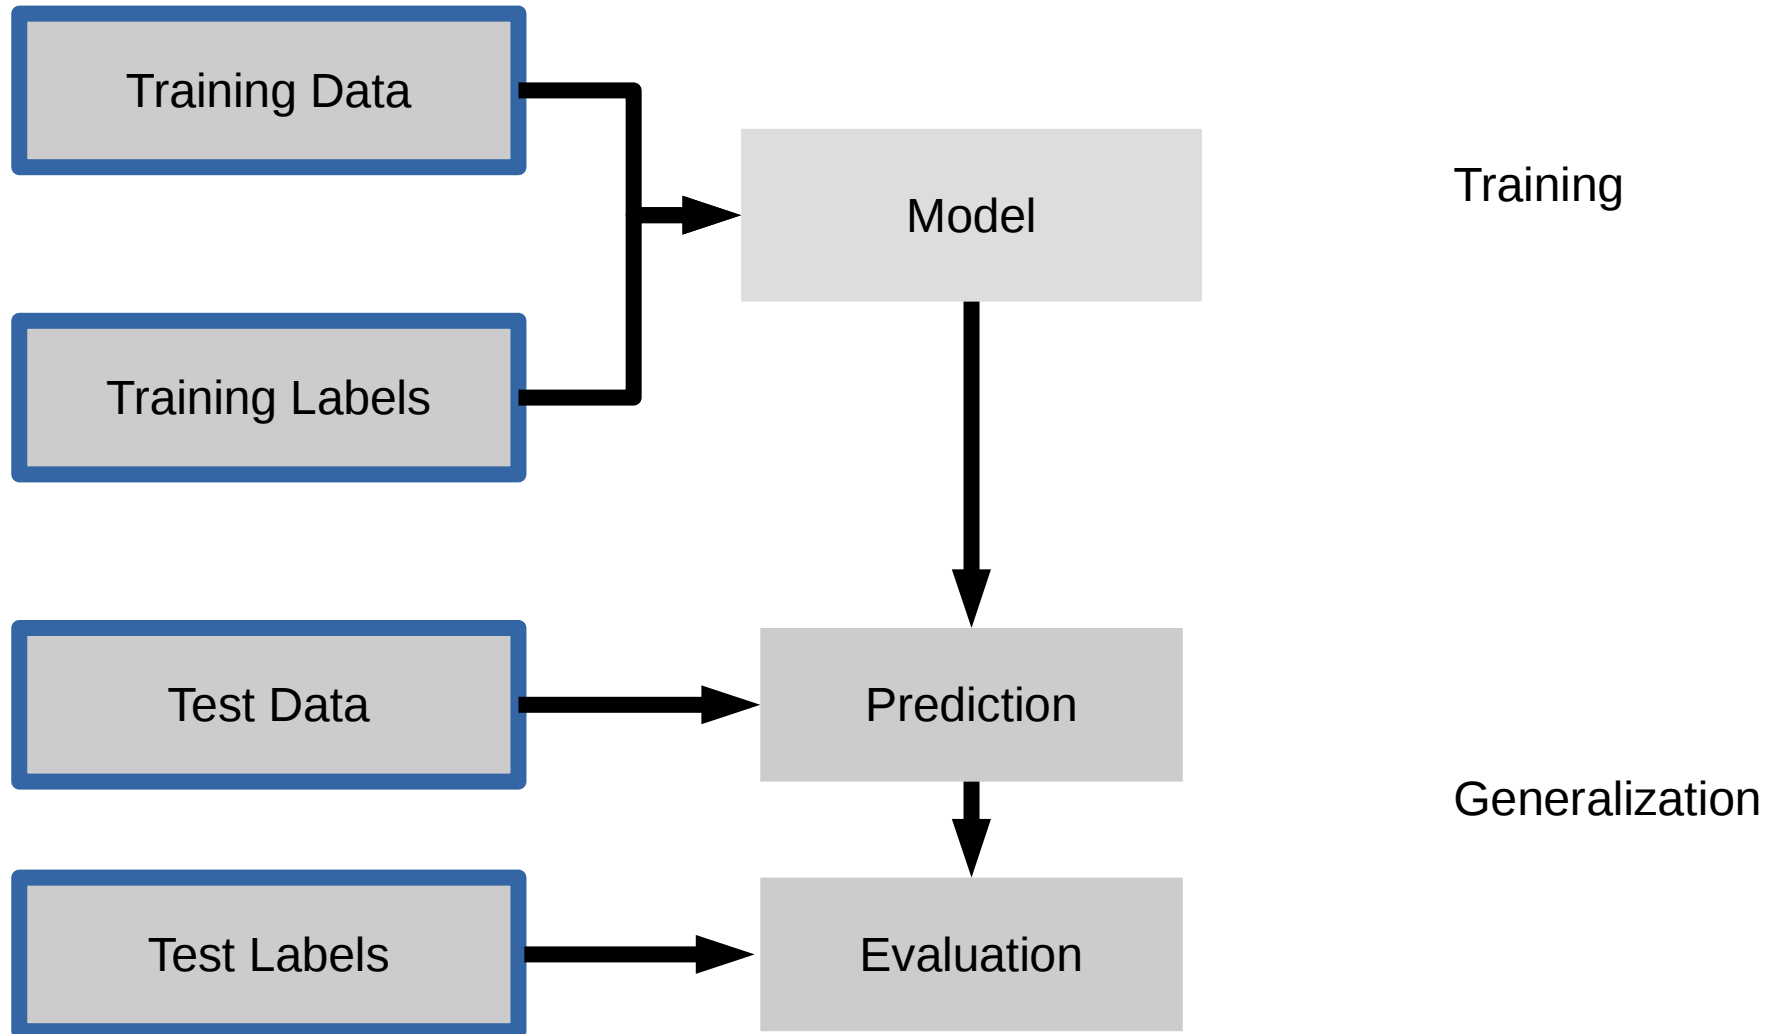
$$X = \begin{pmatrix} 1.1 & 2.2 & 3.4 & 5.6 & 1.0 \\ 6.7 & 0.5 & 0.4 & 2.6 & 1.6 \\ 2.4 & 9.3 & 7.3 & 6.4 & 2.8 \\ 1.5 & 0.0 & 4.3 & 8.3 & 3.4 \\ 0.5 & 3.5 & 8.1 & 3.6 & 4.6 \\ 5.1 & 9.7 & 3.5 & 7.9 & 5.1 \\ 3.7 & 7.8 & 2.6 & 3.2 & 6.3 \end{pmatrix}$$

test set

$$y = \begin{pmatrix} 1.6 \\ 2.7 \\ 4.4 \\ 0.5 \\ 0.2 \\ 5.6 \\ 6.7 \end{pmatrix}$$

# IPython Notebook: Part 0 – Data Loading

# Supervised Machine Learning




```
clf = RandomForestClassifier()
```

```
clf.fit(X_train, y_train)
```

```
y_pred = clf.predict(X_test)
```

```
clf.score(X_test, y_test)
```




# IPython Notebook:

## Part 1 - Introduction to Scikit-learn



# Unsupervised Machine Learning




# Unsupervised Transformations

```
pca = PCA()
```

```
pca.fit(X_train)
```

```
X_new = pca.transform(X_test)
```



# IPython Notebook:

## Part 2 – Unsupervised Transformers

# Basic API

**`estimator.fit(X, [y])`**

**`estimator.predict`**

**`estimator.transform`**

---

Classification

Preprocessing

Regression


Dimensionality reduction

Clustering

Feature selection


Feature extraction

# Nearest neighbors







$$f(x) = y_i, i = \operatorname{argmin}_j ||x_j - x||$$


# Nearest neighbors



# Influence of $n\_neighbors$




# Model Complexity






# Overfitting and Underfitting



# Three-fold split



pro: fast, simple

con: high variance, bad use of data.

```
val_scores = []
neighbors = np.arange(1, 15, 2)
for i in neighbors:
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    val_scores.append(knn.score(X_val, y_val))
print("best validation score: {:.3f}".format(np.max(val_scores)))
best_n_neighbors = neighbors[np.argmax(val_scores)]
print("best n_neighbors: {}".format(best_n_neighbors))


knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(X_trainval, y_trainval)
print("test-set score: {:.3f}".format(knn.score(X_test, y_test)))
```

best validation score: 0.972

best n\_neighbors: 3


test-set score: 0.965

# Cross-validation



Pro: more stable, more data  
con: slower

# Cross-validation + test-set



```

from sklearn.model_selection import cross_val_score

X_train, X_test, y_train, y_test = train_test_split(X, y)

cross_val_scores = []

for i in neighbors:
    knn = KNeighborsClassifier(n_neighbors=i)
    scores = cross_val_score(knn, X_trainval, y_trainval, cv=10)
    cross_val_scores.append(np.mean(scores))

print("best cross-validation score: {:.3f}".format(np.max(cross_val_scores)))
best_n_neighbors = neighbors[np.argmax(cross_val_scores)]
print("best n_neighbors: {}".format(best_n_neighbors))


knn = KNeighborsClassifier(n_neighbors=best_n_neighbors)
knn.fit(X_train, y_train)
print("test-set score: {:.3f}".format(knn.score(X_test, y_test)))

```

```

best cross-validation score: 0.972
best n_neighbors: 3
test-set score: 0.972

```




# IPython Notebook:

## Part 3 – Cross-validation and grid-search




# Preprocessing



```

: plt.boxplot(X)
: plt.xticks(np.arange(1, X.shape[1] + 1), boston.feature_names, rotation=30, ha="right")
: plt.ylabel("MEDV")
: <matplotlib.text.Text at 0x7f580303eac8>

```






```
from sklearn.linear_model import Ridge
X, y = boston.data, boston.target
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)

scaler = StandardScaler()
scaler.fit(X_train)
X_train_scaled = scaler.transform(X_train)
ridge = Ridge().fit(X_train_scaled, y_train)

X_test_scaled = scaler.transform(X_test)
ridge.score(X_test_scaled, y_test)
```

0.63448846877867426



# Categorical Features

# Categorical Features

$\{\text{"red"}, \text{"green"}, \text{"blue"}\} \subset \mathbb{R}^p$  ?

# Categorical Variables


|  | “red” | “green” | “blue” |
|--|-------|---------|--------|
|  | 1     | 0       | 0      |
|  | 0     | 1       | 0      |
|  | 0     | 0       | 1      |



# IPython Notebook: Part 4 – Preprocessing

# Linear Models for Regression

# Linear Models for Regression



$$\hat{y} = w^T \mathbf{x} + b = \sum_{i=1}^p w_i x_i + b$$

# Linear Regression

## Ordinary Least Squares

$$\hat{y} = w^T \mathbf{x} + b = \sum_{i=1}^p w_i x_i + b$$

$$\min_{w \in \mathbb{R}^p} \sum_{i=1}^p ||w^T \mathbf{x}_i - y_i||^2$$

Unique solution if  $\mathbf{X} = (\mathbf{x}_1, \dots, \mathbf{x}_n)^T$  has full rank.

# Ridge Regression

$$\min_{w \in \mathbb{R}^p} \sum_{i=1}^n ||w^T x_i - y_i||^2 + \alpha ||w||^2$$

Always has a unique solution.

Has tuning parameter alpha


# IPython Notebook:

## Part 5 – Linear Models for Regression

# Linear Models for Classification

# Linear models for **binary** classification






$$\hat{y} = \text{sign}(w^T \mathbf{x} + b) = \text{sign}\left(\sum_i w_i x_i + b\right)$$

# Logistic Regression

$$\min_{w \in \mathbb{R}^p} - \sum_{i=1}^n \log(\exp(-y_i w^T \mathbf{x}_i) + 1)$$

$$p(y|\mathbf{x}) = \frac{1}{1 + e^{-w^T \mathbf{x}}}$$




$$\hat{y} = \text{sign}(w^T \mathbf{x} + b)$$


# Penalized Logistic Regression

$$\min_{w \in \mathbb{R}^p} -C \sum_{i=1}^n \log(\exp(-y_i w^T \mathbf{x}_i) + 1) + ||w||_2^2$$


C is inverse to alpha (or alpha / n\_samples)




$C = 0.010000$



$C = 10.000000$



$C = 1000.000000$



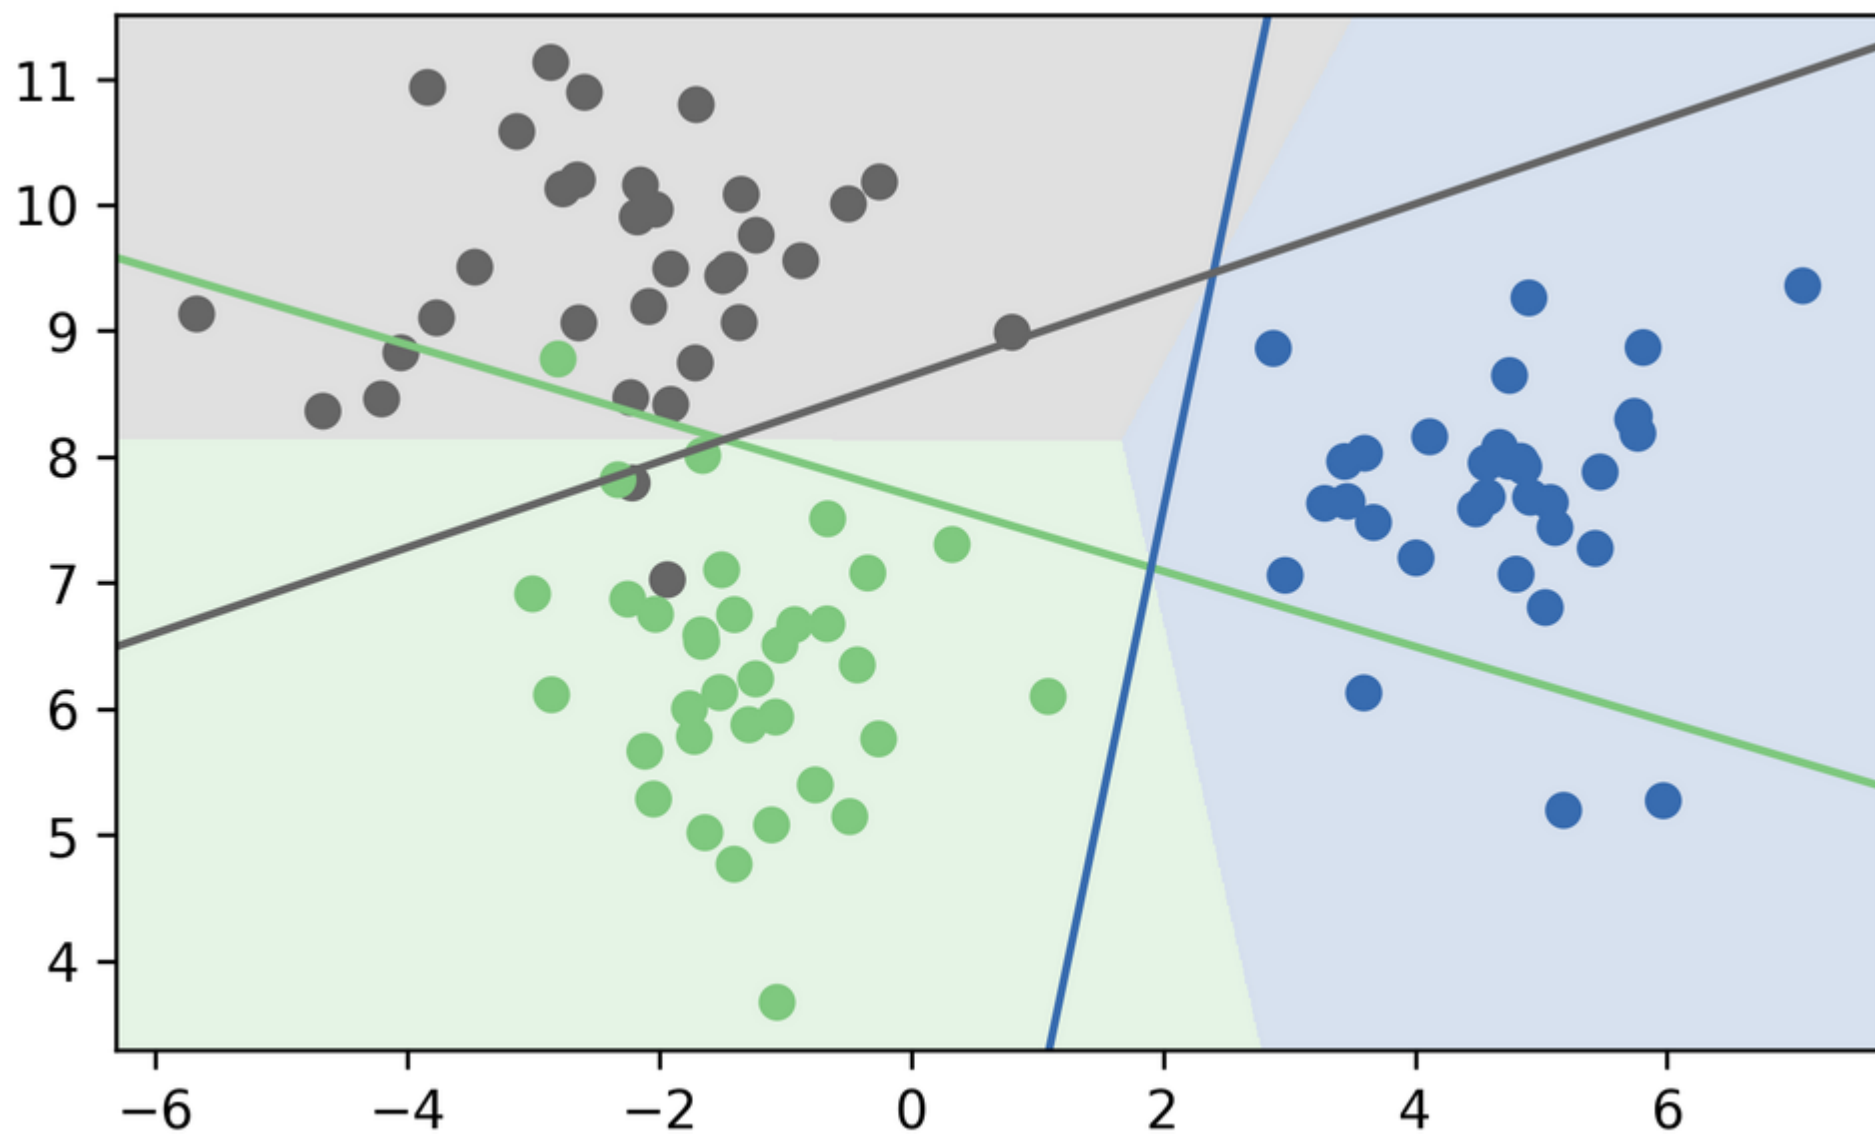
# Multinomial Logistic Regression

Probabilistic multi-class model:

$$p(y = i | \mathbf{x}) = \frac{e^{-\mathbf{w}_i^T \mathbf{x}}}{\sum_{j \in Y} e^{-\mathbf{w}_j^T \mathbf{x}}}$$

$$\min_{w \in \mathbb{R}^p} - \sum_{i=1}^n \log(p(y = y_i | \mathbf{x}_i))$$

$$\hat{y} = \arg \max_{i \in Y} \mathbf{w}_i \mathbf{x}$$



$$\hat{y} = \arg \max_{i \in Y} \mathbf{w}_i \mathbf{x}$$

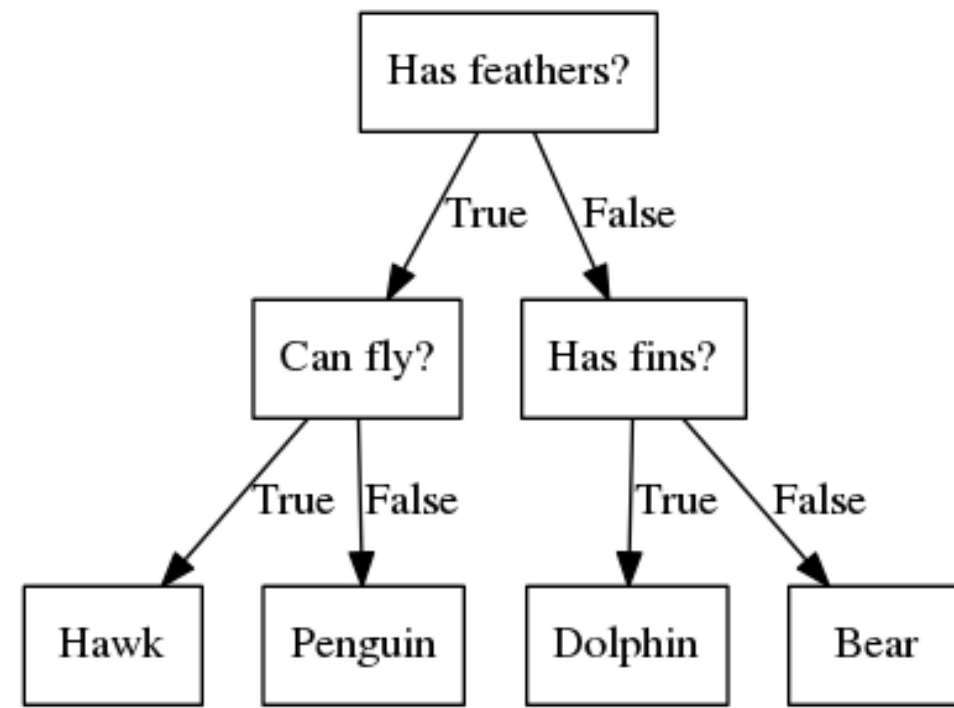
# IPython Notebook:

## Part 6 – Linear Models for Classification

# Decision Trees and Tree-based Models





# Idea: series of binary questions





# Building trees

depth = 1



depth = 2




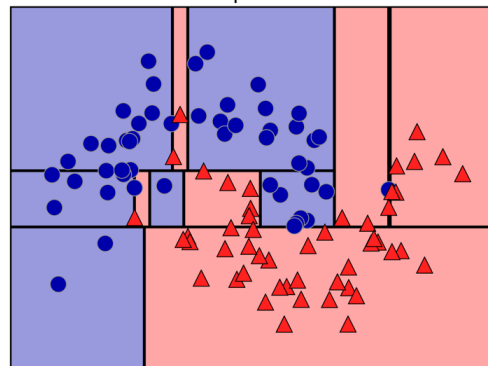
Continuous features:  
“questions” are thresholds on  
single features.

[Other methods are possible  
but not as common]


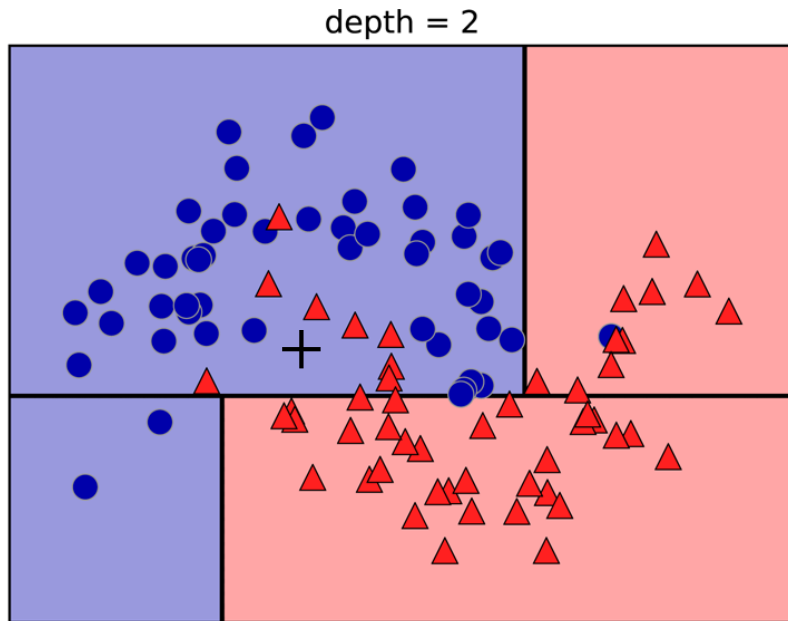
For each split:  
exhaustive search over all  
features and thresholds!

Minimize “impurity”

depth = 9



# Prediction

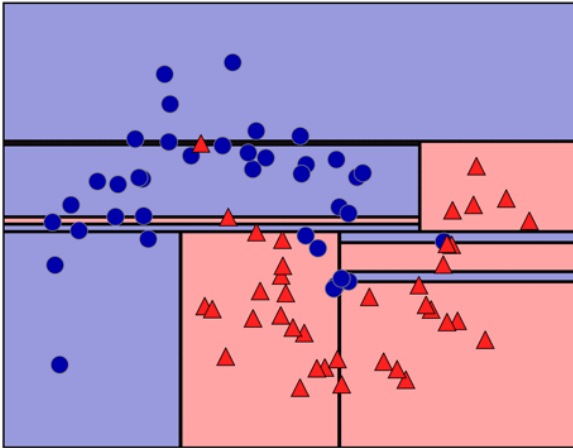


Traverse tree based on feature tests  
Predict most common class in leaf

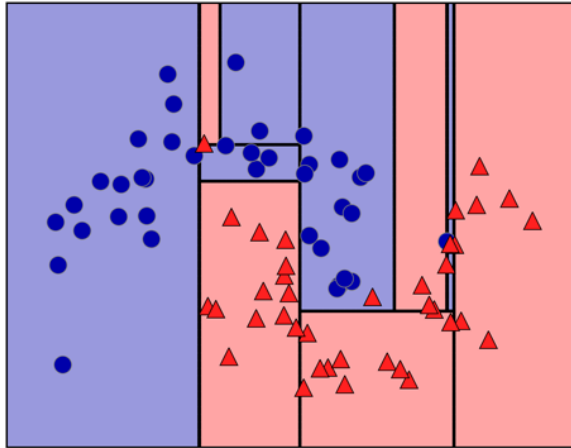
# IPython Notebook: Part 7 – Decision Trees

# Random Forests

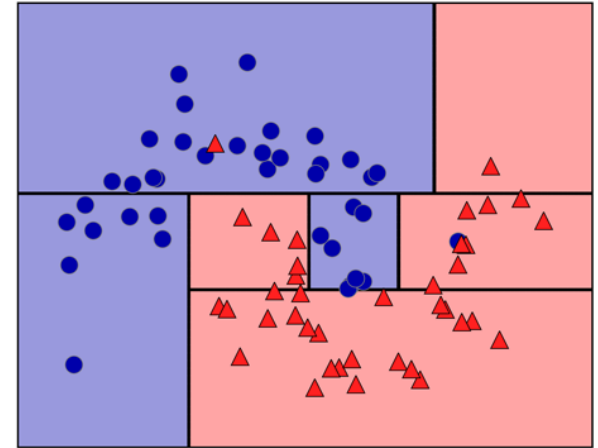
tree 0




tree 1




tree 2




tree 3



tree 4

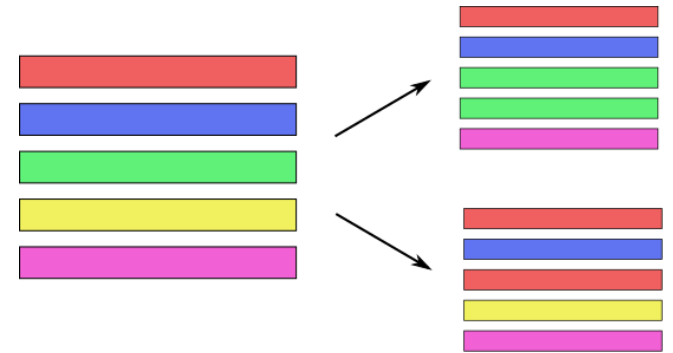


random forest

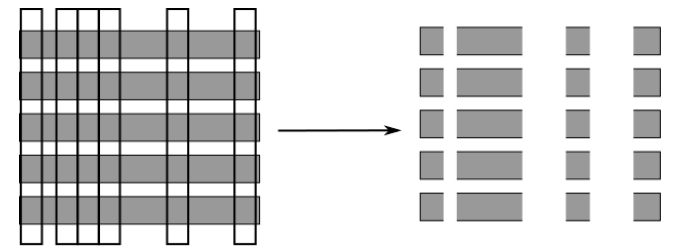


# Randomize in two ways

- For each **tree**:  
Pick bootstrap sample of data



- For each **split**:  
Pick random sample of features




- More tree are always better

# Tuning Random Forests

- Main parameter: `max_features`
  - around `sqrt(n_features)` for classification
  - Around `n_features` for regression
- `n_estimators > 100`
- Prepruning might help, definitely helps with model size!
- `max_depth`, `max_leaf_nodes`, `min_samples_split` again

O'REILLY®



# Introduction to Machine Learning with Python

A GUIDE FOR DATA SCIENTISTS

Andreas C. Müller & Sarah Guido




[amueller.github.io](https://amueller.github.io)



[@amuellerm1](https://twitter.com/amuellerm1)



[@amueller](https://github.com/amueller)



[andreas.mueller@columbia.edu](mailto:andreas.mueller@columbia.edu)